

Parallel & Distributed Computer Systems 2023

Homework 2

Distributed k-select

Deadline: January 13, 2024

Introduction

As we discussed in class, the MPI homework is the distributed k-select using MPI. First, we write a sequential function to do partition in place, as in this MATLAB.

```
function [p,q,A] = partitioninplace(A,v)
% function [p,q,A] = partitioninplace(A,v)
% reorder array A(1:n) and find p and q so that after the call,
% A(1:p-1) < v, A(p:q-1) == v and A(q:n) > v
% for any value v that may or may not be in A(1:n)
```

Use the above function to implement `kselect(A,k)` that finds the kth smallest value of array `A`. Integer `k` is in the range `1:n` and `A` is an array of length `n`. The function should return the value `A(k)` of a sorted `A`, without sorting `A`.

Parallel implementation (C/C++)

Once your model code in MATLAB/Python/Julia/R (whatever allows you to quickly prototype and confirm correctness) works, implement it in C/C++.

Suppose we had a very large file holding the array `A(1:n)` (that exceeds the memory of a computer). Extend your sequential code with MPI into a distributed program for `kselect(A,k)`, where `p` is the number of processes participating in the computation. Each process `1:p` works with its own data, but it coordinates globally for every decision how to proceed with the search. It is not worth to redistribute the data unless `n/p` fits in a single memory.

Extra credit if you use multiple threads in each process to parallelize `partition` and take advantage of the shared memory of a multicore processor. A multi-threaded solution will need to be using a ping-pong buffer as I think it will be too hard to do it in place and in parallel and faster :smile:. Feel free to use any threading compiler or library but make sure it is compatible with the MPI library you are using.

Alternative Parallel implementation in Julia

If you elect to implement the whole project in Julia and not C/C++, the multithreading won't be optional. Let's discuss what packages to use. If you'd like to use some other language, please coordinate with me first.

What to submit

- A 4 page report in PDF format. Report execution times and speedup of your implementations with respect to the number of items n , and the number of processes p . Please identify how many different host servers are used. MPI behaves differently if the processes are on the same host or on different hosts.
- Upload the source code on GitHub, BitBucket, Dropbox, Google Drive, etc. and add a link in your report.
- Argue about the validity and effectiveness of your code. You can also include any other information you consider important for the evaluation of your work.
- Cite any external sources you may have used.